

An introduction to Bayesian modelling with R, JAGS and STAN

Francisco Rodriguez-Sanchez (@frod_san)

November 2016

Ecology Letters, (2005) **8**: 2–14

doi: 10.1111/j.1461-0248.2004.00702.x

IDEAS AND PERSPECTIVES

Why environmental scientists are becoming Bayesians

- Powerful
- Flexible
- Knowledge synthesis
- Uncertainty

Even if you don't become Bayesian today,
this workshop will help you to understand
regression and 'mixed-effects' models better.

*The purpose of models is not to fit data,
but to sharpen the questions*

Samuel Karlin

General

- Data analysis using regression and multilevel/hierarchical models
- Statistical Rethinking: A Bayesian Course with Examples in R and Stan
- Bayesian data analysis
- The BUGS book
- Doing Bayesian data analysis: a tutorial with R, JAGS, and STAN
- Bayesian linear mixed models using Stan: a tutorial
- Bayesian basics
- Philosophy and the practice of Bayesian statistics

Ecology-oriented

- The ecological detective: confronting models with data
- Bayesian methods for ecology
- Models for ecological data
- Introduction to WinBUGS for ecologists
- Applied hierarchical modeling in ecology
- Bayesian Models: A Statistical Primer for Ecologists
- and more. . .

- WinBUGS/OpenBUGS
- JAGS
- STAN
- Filzbach
- Nimble
- Many R packages: MCMCpack, MCMCgImm, LaplacesDemon, r-inla, etc
(see [Bayesian task view](#))

We'll focus on JAGS and STAN

- Fast, powerful, and most popular
- Similar to BUGS
- Easy to start
- Open-ended modelling: deal with complex models too
- But look for specific implementations of your analysis (e.g. hSDM, rstanarm)
- Once concepts understood, switching software not difficult

JAGS has to be installed independently

<http://mcmc-jags.sourceforge.net/>

Use latest version (4.2.0)

- `rjags`
- `R2jags`
- `runjags`
- `jagsUI`
- `dclone`
- `rube`

- `rstan`
- `rstanarm`
- `rethinking`
- `brms`


We will also need these R packages:

```
library(arm)
library(R2jags)
library(ggmcmc)
library(shinystan)
library(rube)
library(lme4)
library(rstan)
library(rstanarm)
```

and their dependencies

The very basics: linear regression


The very basics: linear regression



$$y_i = a + bx_i + \epsilon_i$$

How many parameters?

The very basics: linear regression



Or also

$$y_i = a + bx_i + \epsilon_i$$

$$\epsilon_i \sim N(0, \sigma^2)$$

$$y_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = a + bx_i$$

Our dataset: tree heights and DBH


<http://tinyurl.com/treesdata>

- One species
- 10 plots
- 1000 trees
- Number of trees per plot ranging from 4 to 392

```
trees <- read.csv("trees.csv")  
summary(trees[, 1:3])
```

	plot	dbh	height
Min.	: 1.0	Min. : 5.06	Min. :13.40
1st Qu.:	1.0	1st Qu.:17.69	1st Qu.:29.68
Median :	2.0	Median :28.62	Median :36.55
Mean :	2.7	Mean :27.88	Mean :36.51
3rd Qu.:	4.0	3rd Qu.:38.97	3rd Qu.:43.33
Max.	:10.0	Max. :49.92	Max. :59.30

What's the relationship between DBH and height?



First step: linear regression (lm)

```
simple.lm <- lm(height ~ dbh, data = trees)
arm::display(simple.lm) # summary of key model elements
```

```
lm(formula = height ~ dbh, data = trees)
```

	coef.est	coef.se
(Intercept)	19.34	0.31
dbh	0.62	0.01

```
n = 1000, k = 2
```

```
residual sd = 4.09, R-Squared = 0.79
```

Interpretation?


```
summary(trees$dbh)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.06	17.69	28.62	27.88	38.96	49.92

```
trees$dbh.c <- trees$dbh - 25
```

So, all parameters will be referred to a 25 cm DBH tree.

Linear regression with centred DBH



```
lm(formula = height ~ dbh.c, data = trees)
      coef.est coef.se
(Intercept)  34.73    0.13
dbh.c         0.62    0.01
---
n = 1000, k = 2
residual sd = 4.09, R-Squared = 0.79
```

Things we'll need

- A function describing the model (including **priors**)
- Data
- Choose parameters to save
- Define initial values for MCMC chains
- Decide number of iterations (and burnin)

$$y_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \alpha + \beta x_i$$

In this case:

$$\text{Height}_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \alpha + \beta \text{DBH}_i$$

α : expected height when $\text{DBH} = 0$

β : how much height increases with every unit increase of DBH

$$\tau = \frac{1}{\sigma^2}$$

So, residual variance $\sigma^2 = 100$ expressed as $\tau = 0.01$.

Specify the model as an R function

$$\text{Height}_i \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \alpha + \beta \text{DBH}_i$$

```
bayes.lm <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau=precision (inverse var)  
    mu[i] <- alpha + beta*dbh[i]      # expected height ~ dbhc  
  }  
  
}
```


We need priors for every parameter!

```
bayes.lm <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)    # tau = precision (inverse var)  
    mu[i] <- alpha + beta*dbh[i]    # expected height ~ dbhc  
  }  
  
  # PRIORS (vague or weakly informative)  
  alpha ~ dunif(1, 100)    # prior avg height of 25-cm-DBH tree  
  beta ~ dunif(0, 10)    # how much do we expect height to scale with DBH?  
  tau <- 1/(sigma*sigma)    # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)    # residual standard deviation  
}
```

Avoid 'non-informative' priors (see [this](#) and [this](#))

Use *weakly informative* (e.g. bounded Uniform, Normal with reasonable parameters, Cauchy...)

or *strongly informative* priors based on previous knowledge and common sense.

Some tips for setting priors:

- <https://github.com/stan-dev/stan/wiki/Prior-Choice-Recommendations>
- http://www.mrc-bsu.cam.ac.uk/wp-content/uploads/bugsbook_chapter5.pdf


Usually good idea to try different priors and evaluate posterior sensitivity

Or run model without likelihood (priors only).

Example: estimating people height across countries

Unreasonable prior


```
plot(density(rnorm(1000, 0, 1000)),  
     main="", xlab="Height (m)")
```



(from STAN manual)

Reasonable prior

```
plot(density(rnorm(1000, 2, 0.5)),  
     main="", xlab="Height (m)")
```



We already have our model definition

```
bayes.lm <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha + beta*dbh[i]      # centred diameter  
  }  
  
  # PRIORS (vague or weakly informative)  
  alpha ~ dunif(1, 100)               # prior for average height of a 25-cm-DBH t  
  beta ~ dunif(0, 10)                 # how much do we expect height to scale wit  
  tau <- 1/(sigma*sigma)               # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)                # residual standard deviation  
}
```

Next step: create list with data

Data = known values

```
data <- list(height = trees$height,  # response  
             dbhc = trees$dbh.c)    # predictor
```

Next step: choose parameters to save

```
params <- c("alpha", "beta", "sigma")
```

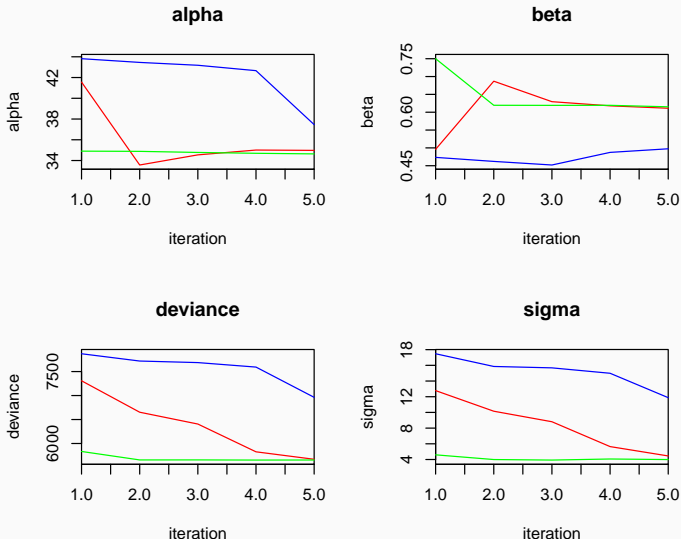
Now call JAGS to run the model

```
library(R2jags)
```

```
m1 <- jags(data,  
  model.file = bayes.lm,  
  parameters.to.save = params,  
  n.chains = 3,  
  inits = NULL,    # JAGS will create inits for each chain  
  n.iter = 10,     # number of iterations  
  n.burnin = 5)    # iterations to discard (before convergence)
```

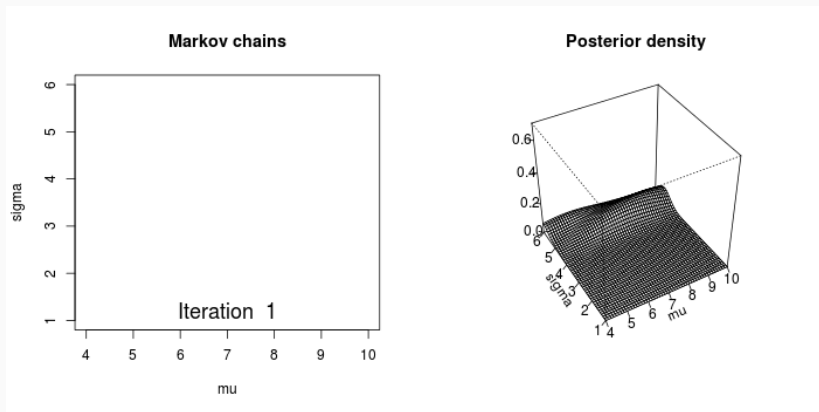
Traceplots: viewing MCMC in action

```
traceplot(m1, ask = FALSE, mfrow = c(2, 2))
```



Obviously we haven't achieved convergence yet

Viewing MCMC in action (II)



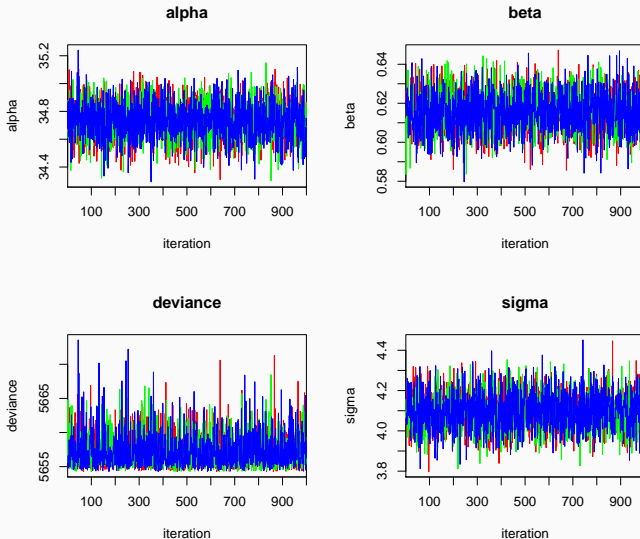
Source: <http://mbjoseph.github.io/2013/09/08/metropolis.html>

Let's run JAGS for longer

```
m1 <- jags(data,  
  model.file = bayes.lm,  
  parameters.to.save = params,  
  n.chains = 3,  
  inits = NULL,  
  n.iter = 10000,      # 10000 MCMC iterations  
  n.burnin = 5000)    # discard first half (5000 iterations)
```

Use traceplots to assess convergence

```
traceplot(m1, ask = FALSE, mfrow = c(2, 2))
```



Results: parameter estimates

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpisLMp5
  3 chains, each with 10000 iterations (first 5000 discarded), n.thin =
  n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alpha	34.732	0.134	34.469	34.992	1.001	3000
beta	0.616	0.010	0.596	0.635	1.002	1500
sigma	4.095	0.090	3.924	4.276	1.001	2000
deviance	5657.236	2.422	5654.464	5663.367	1.006	770

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)

$pD = 2.9$ and $DIC = 5660.2$

DIC is an estimate of expected predictive error (lower deviance is better)



```
lm(formula = height ~ dbh.c, data = trees)

            coef.est coef.se
(Intercept) 34.73      0.13
dbh.c        0.62      0.01
---
n = 1000, k = 2
residual sd = 4.09, R-Squared = 0.79
```

Results pretty similar, because of vague priors

A plot of the whole model

Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpisLMp5/model27982df47934.txt", fit using jags, 3 chains, each with 10000 iterations (first 5000 discarded)



Model checking

Using ggcmc to produce diagnostic plots

```
suppressPackageStartupMessages(library(ggcmc))  
m1.mcmc <- as.mcmc(m1)      # Get list of MCMC values  
m1.tidy = ggs(m1.mcmc)      # Produce tidy data frame  
ggcmc(m1.tidy)
```

Plotting histograms

Plotting density plots

Plotting traceplots

Plotting running means

Plotting comparison of partial and full chain

Plotting autocorrelation plots

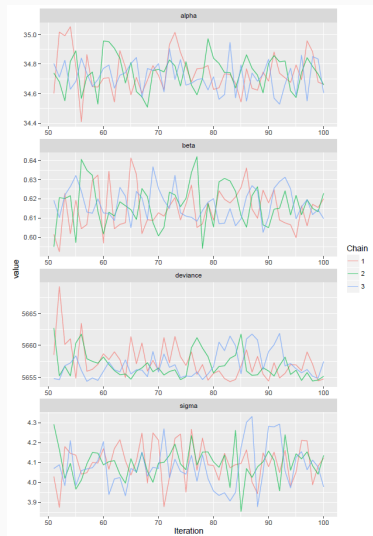

Plotting crosscorrelation plot

Plotting Potential Scale Reduction Factors

Plotting Geweke Diagnostic

Plotting caterpillar plot


ggmcmc output sample



```
library(shinystan)  
launch_shinystan(as.shinystan(m1.mcmc))
```


Comparing prior and posterior


```
library(rube)
priPost(post = m1$BUGSoutput$sims.list$alpha[, 1],
        dist = "Uniform", pripar = c(1, 100))
```




Comparing prior and posterior (sigma)

```
priPost(post = m1$BUGSoutput$sims.list$sigma[, 1],  
        dist = "Uniform", pripar = c(0, 50))
```

Prior / Posterior Plot




Observed vs Predicted values




How would you do this?

Observed vs Predicted values

```
alpha.avg <- mean(m1$BUGSoutput$sims.list$alpha)
beta.avg <- mean(m1$BUGSoutput$sims.list$beta)
mu <- alpha.avg + beta.avg*trees$dbh.c
plot(trees$height, mu,
     xlim = c(0, 60), ylim = c(0, 60),
     xlab = "Observed height (m)", ylab = "Predicted height (m)")
abline(a = 0, b = 1)
```





Now using Normal vague priors

Model with Normal priors

```
bayes.lm.N <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha + beta*dbhc[i]      # centred diameter  
  }  
  
  # PRIORS  
  alpha ~ dnorm(0, 0.01)              # prior for intercept  
  beta ~ dnorm(0, 0.01)               # prior for beta (slope)  
  tau <- 1/(sigma*sigma)              # tau = 1/sigma2  
  sigma ~ dunif(0, 50)               # residual standard deviation  
}
```

```
m1b <- jags(data,  
            model.file = bayes.lm.N,  
            parameters.to.save = params,  
            n.chains = 3,  
            inits = NULL,  
            n.iter = 4000,  
            n.burnin = 2000)
```

Results with Normal priors on intercept and slope

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpisLmp5
  3 chains, each with 4000 iterations (first 2000 discarded), n.thin = 2
n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alpha	34.727	0.130	34.466	34.978	1.001	3000
beta	0.616	0.010	0.597	0.635	1.001	3000
sigma	4.095	0.091	3.926	4.272	1.001	2000
deviance	5657.132	2.368	5654.429	5663.500	1.001	3000

For each parameter, n.eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor (at convergence, Rhat=

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)

$pD = 2.8$ and $DIC = 5659.9$

DIC is an estimate of expected predictive error (lower deviance is better)

Very similar

Let's try strongly informative prior on alpha

```
bayes.lm.N2 <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha + beta*dbh[i]      # centred diameter  
  }  
  
  # PRIORS  
  alpha ~ dnorm(10, 10)               # prior for intercept  
  beta ~ dnorm(0, 0.01)               # prior for beta (slope)  
  tau <- 1/(sigma*sigma)              # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)               # residual standard deviation  
}
```

```
m1c <- jags(data,  
            model.file = bayes.lm.N2,  
            parameters.to.save = params,  
            n.chains = 3,  
            inits = NULL,  
            n.iter = 4000,  
            n.burnin = 2000)
```

Posteriors: between prior and likelihood

```
Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpisLMp5
  3 chains, each with 4000 iterations (first 2000 discarded), n.thin = 2
n.sims = 3000 iterations saved
```

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alpha	14.801	0.352	14.104	15.495	1.001	2800
beta	0.950	0.048	0.856	1.042	1.001	3000
sigma	19.900	0.553	18.831	20.998	1.001	3000
deviance	8817.718	33.837	8750.531	8884.142	1.001	2900

For each parameter, n.eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor (at convergence, Rhat=

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)


$pD = 572.5$ and $DIC = 9390.2$

DIC is an estimate of expected predictive error (lower deviance is better)

Prior vs Posterior

```
priPost(post = m1c$BUGSoutput$sims.list$alpha[,1],  
        dist = "Normal", pripar = c(10, 0.1))
```

Prior / Posterior Plot



Bayesian inference

$$P(\text{Hypothesis}|\text{Data}) = \frac{P(\text{Data}|\text{Hypothesis})P(\text{Hypothesis})}{P(\text{Data})}$$


$$P(\text{Hypothesis}|\text{Data}) \propto P(\text{Data}|\text{Hypothesis}) \times P(\text{Hypothesis})$$

$$\text{Posterior} \propto \text{Likelihood} \times \text{Prior}$$

Hence, prior influence decreases with size of data set.

Varying-intercept models

Accounting for plot effects



Do it yourself using `lm`

lm results

```
lm.plot <- lm(height ~ factor(plot) + dbh.c, data = trees)
```

```
lm(formula = height ~ factor(plot) + dbh.c, data = trees)
```

	coef.est	coef.se
(Intercept)	32.13	0.16
factor(plot)2	6.50	0.26
factor(plot)3	4.36	0.35
factor(plot)4	1.93	0.36
factor(plot)5	3.64	0.34
factor(plot)6	4.20	0.42
factor(plot)7	-0.18	0.67
factor(plot)8	-5.31	0.89
factor(plot)9	5.44	1.09
factor(plot)10	2.26	1.37
dbh.c	0.62	0.01


```
n = 1000, k = 11
```

```
residual sd = 3.04, R-Squared = 0.88
```


Interpretation?

Single vs varying intercept

Pooling all plots



Different intercept for each plot



Things we'll need

- A function describing the model (including **priors**)
- Data
- Choose parameters to save
- Define initial values for MCMC chains
- Decide number of iterations (and burnin)

Bayesian varying-intercept model with no pooling

```
varint.nopool <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha[plot[i]] + beta*dbhc[i]    # centred diameter  
  }  
  # PRIORS  
  #alpha ~ dnorm(0, .001)      # previous model  
  for (j in 1:10){  
    alpha[j] ~ dnorm(0, .001) # Plot effects drawn from Normal distrib  
                                # with large **fixed** variance  
  }  
  beta ~ dnorm(0, .001)  
  tau <- 1/(sigma*sigma)      # tau = 1/sigma^2  
  sigma ~ dunif(0, 50)  
}
```

This fits same model as `lm.plot`

```
data <- list(height = trees$height,  
             dbhc = trees$dbh.c,  
             plot = trees$plot)  
m2 <- jags(data,  
           model.file = varint.nopool,  
           parameters.to.save = params,  
           n.chains = 3,  
           inits = NULL,  
           n.iter = 4000,  
           n.burnin = 2000)
```

Running JAGS in parallel (multiple cores)

```
m2 <- jags.parallel(data,  
  model.file = varint.nopool,  
  parameters.to.save = params,  
  n.chains = 3,  
  inits = NULL,  
  n.iter = 4000,  
  n.burnin = 2000)
```

Results of Bayesian varying intercept model without pooling

Inference for Bugs model at "varint.nopool", fit using jags,
3 chains, each with 4000 iterations (first 2000 discarded), n.thin = 2
n.sims = 3000 iterations saved

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alpha[1]	32.126	0.154	31.821	32.427	1.001	3000
alpha[2]	38.631	0.202	38.228	39.033	1.002	3000
alpha[3]	36.487	0.328	35.830	37.109	1.002	1700
alpha[4]	34.050	0.322	33.397	34.663	1.001	2200
alpha[5]	35.754	0.309	35.149	36.346	1.002	1400
alpha[6]	36.339	0.389	35.550	37.094	1.002	1100
alpha[7]	31.934	0.632	30.746	33.168	1.001	3000
alpha[8]	26.810	0.870	25.174	28.493	1.001	3000
alpha[9]	37.519	1.089	35.460	39.709	1.001	3000
alpha[10]	34.341	1.367	31.708	36.983	1.001	3000
beta	0.617	0.008	0.603	0.631	1.003	950
sigma	3.047	0.067	2.919	3.183	1.003	680
deviance	5064.489	4.777	5056.981	5075.523	1.002	1000

For each parameter, n.eff is a crude measure of effective sample size,⁶⁸

Same results as `lm.plot`

```
lm(formula = height ~ factor(plot) + dbh.c, data = trees)
```




	coef.est	coef.se
(Intercept)	32.13	0.16
factor(plot)2	6.50	0.26
factor(plot)3	4.36	0.35
factor(plot)4	1.93	0.36
factor(plot)5	3.64	0.34
factor(plot)6	4.20	0.42
factor(plot)7	-0.18	0.67
factor(plot)8	-5.31	0.89
factor(plot)9	5.44	1.09
factor(plot)10	2.26	1.37
dbh.c	0.62	0.01

```
n = 1000, k = 11
```

```
residual sd = 3.04, R-Squared = 0.88
```

Plot whole model

Bugs model at "varint.nopool", fit using jags, 3 chains, each with 4000 iterations (first 2000 discarded)



The varying-intercept model seems better

Based on Deviance Information Criterion (DIC)

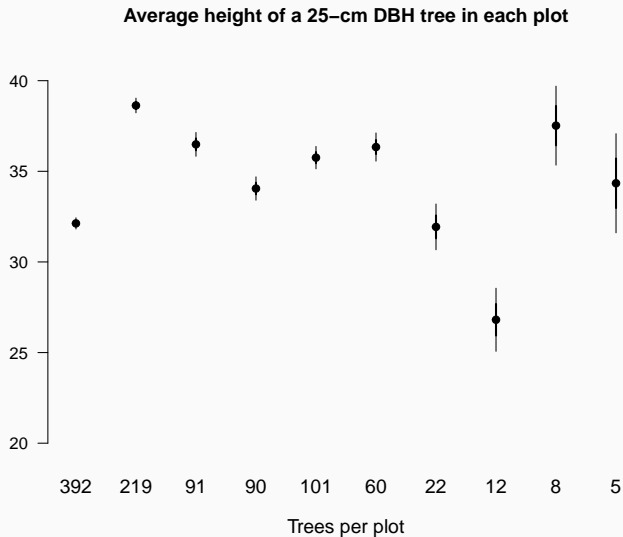
$$\text{DIC}(m1) = 5660$$

$$\text{DIC}(m2) = 5076$$

But note DIC (as AIC and BIC) have been criticised

WAIC and cross-validation are preferred alternatives.

Estimation of plot effects improves with sample size



Varying-intercepts with pooling (shrinkage)

$$y_i = a_j + bx_i + \varepsilon_i$$

$$a_j \sim N(0, \tau^2)$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

In our example:

$$Height_i = plot_j + bDBH_i + \varepsilon_i$$

$$plot_j \sim N(0, \tau^2)$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

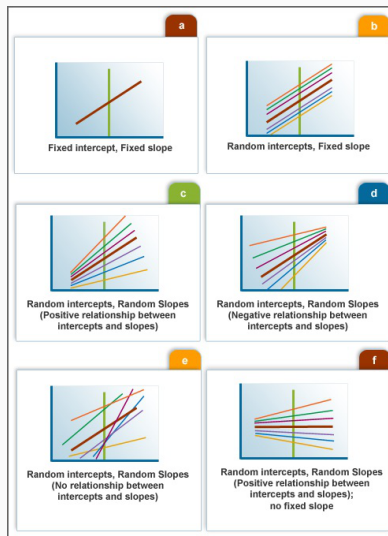
Many definitions, so much confusion

What is a random effect, really?

1. Varies by group
2. Variation estimated with probability model

Random effects are estimated with partial pooling, while fixed effects are not (infinite variance).

Random vs fixed effects



<http://www.esourceresearch.org/>

Fitting mixed models with lmer

```
mixed <- lmer(height ~ dbh.c + (1|plot), data = trees)
```

```
lmer(formula = height ~ dbh.c + (1 | plot), data = trees)
```

	coef.est	coef.se
(Intercept)	34.43	1.08
dbh.c	0.62	0.01

Error terms:

Groups	Name	Std.Dev.
plot	(Intercept)	3.35
Residual		3.04

number of obs: 1000, groups: plot, 10

AIC = 5116.3, DIC = 5096.3

deviance = 5102.3

lmer coefficients

```
coef(mixed)
```

```
$plot
```

	(Intercept)	dbh.c
1	32.13118	0.6169271
2	38.61479	0.6169271
3	36.46547	0.6169271
4	34.06404	0.6169271
5	35.75313	0.6169271
6	36.30517	0.6169271
7	32.04003	0.6169271
8	27.30620	0.6169271
9	37.27097	0.6169271
10	34.39546	0.6169271

```
attr("class")
```

```
[1] "coef.mer"
```

Bayesian varying-intercept model with pooling across plots

```
varint.pool <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha[plot[i]] + beta*dbh[i]    # centred diameter  
  }  
  # PRIORS  
  for (j in 1:10){  
    alpha[j] ~ dnorm(grandmu, tauplot) # Now we are estimating the plo  
  }  
  grandmu ~ dnorm(0, .001)      # Overall mean height across all plots  
  tauplot <- 1/(sigmaplot*sigmaplot)  
  sigmaplot ~ dunif(0, 20)      # between-plot variance  
  beta ~ dnorm(0, .001)  
  tau <- 1/(sigma*sigma)  
  sigma ~ dunif(0, 50)          # residual variance  
}
```

```
data <- list(height = trees$height,  
             dbhc = trees$dbh.c,  
             plot = trees$plot)  
m3 <- jags.parallel(data,  
                    model.file = varint.pool,  
                    parameters.to.save = c("alpha", "beta", "sigma", "grandmu",  
                    n.chains = 3,  
                    inits = NULL,  
                    n.iter = 4000,  
                    n.burnin = 2000)
```



Results of multilevel model with pooling

Inference for Bugs model at "varint.pool", fit using jags,
3 chains, each with 4000 iterations (first 2000 discarded), n.thin = 2
n.sims = 3000 iterations saved


	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
alpha[1]	32.133	0.156	31.827	32.440	1.001	3000
alpha[2]	38.616	0.205	38.211	39.010	1.001	2800
alpha[3]	36.455	0.317	35.839	37.090	1.002	1400
alpha[4]	34.067	0.318	33.436	34.689	1.002	1700
alpha[5]	35.760	0.303	35.174	36.352	1.001	3000
alpha[6]	36.308	0.389	35.544	37.093	1.002	1800
alpha[7]	32.022	0.636	30.818	33.296	1.001	3000
alpha[8]	27.273	0.884	25.561	28.951	1.001	3000
alpha[9]	37.299	1.056	35.230	39.347	1.001	2700
alpha[10]	34.414	1.290	31.822	36.969	1.001	3000
beta	0.617	0.007	0.602	0.631	1.002	1900
grandmu	34.372	1.280	31.696	36.863	1.001	2500
sigma	3.044	0.070	2.913	3.184	1.001	3000
sigmaplot	3.927	1.162	2.353	6.881	1.001	3000
deviance	5064.690	4.881	5056.896	5075.604	1.002	1300

A plot of the whole model

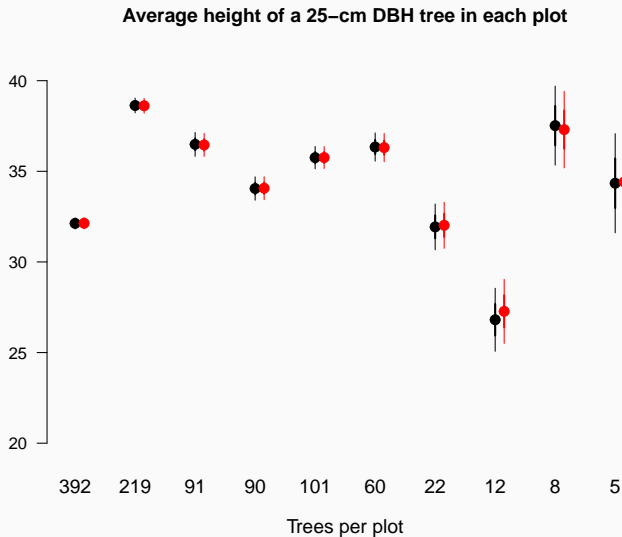
Bugs model at "varint.pool", fit using jags, 3 chains, each with 4000 iterations (first 2000 discarded)



medians and 80% intervals




Comparing plot coefficients




A gradient from complete to no pooling

The multilevel model (with pooling) is somewhere between the complete-pooling (single intercept) and the no-pooling (one intercept for each plot, without shrinkage) models.

Pooling all plots



Different intercept for each plot



Growing the hierarchy: adding plot-level predictors

Model with group-level predictors

We had:

$$y_i = a_j + bx_i + \varepsilon_i$$

$$a_j \sim N(0, \tau^2)$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

Now

$$y_i = a_j + bx_i + \varepsilon_i$$

$$a_j \sim N(\mu_j, \tau^2)$$

$$\mu_j = \gamma + \delta \cdot \text{predictor}_j$$

$$\varepsilon_i \sim N(0, \sigma^2)$$

```
plotdata <- read.csv("plotdata.csv")  
temp.c <- plotdata$temp - 15
```


Model with group-level predictors

```
group.preds <- function(){  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)  
    mu[i] <- alpha[plot[i]] + beta*dbh[i]  
  }  
  # PRIORS  
  for (j in 1:10){  
    alpha[j] ~ dnorm(grandmu + beta.temp*tempc[j], tauplot)  
  }  
  beta.temp ~ dnorm(0, .001)  # slope for temperature effects  
  grandmu ~ dnorm(0, .001)  
  tauplot <- 1/(sigmaplot*sigmaplot)  
  sigmaplot ~ dunif(0, 20)  
  beta ~ dnorm(0, .001)  
  tau <- 1/(sigma*sigma)  
  sigma ~ dunif(0, 50)  
}
```




```
data <- list(height = trees$height,  
             dbhc = trees$dbh.c,  
             plot = trees$plot,  
             tempc = temp.c)  
params = c("alpha","beta","sigma", "grandmu", "sigmaplot", "beta.temp")  
m4 <- jags.parallel(data,  
                    model.file = group.preds,  
                    parameters.to.save = params,  
                    n.chains = 3,  
                    inits = NULL,  
                    n.iter = 4000,  
                    n.burnin = 2000)
```


Average heights among plots related to temperature



Adding plot-level predictors (pooling) may improve parameter estimation



Adding plot-level predictors (pooling) may improve parameter estimation



From Gelman & Hill p. 253

- and coefficients be estimated with pooling
- but taking into account correlation between slopes and intercepts
- (see e.g. Gelman & Hill 2007, ch. 13)

Recapitulating...

So what's a multilevel/hierarchical model?

- Parameters/coefficients are given a probability model (with their own hyperparameters estimated from data).
- Intercepts and/or slopes may vary, and can be modelled (sometimes including their own predictors).

Advantages of hierarchical Bayes


- Perfect for structured data (space-time)
- Predictors enter at the appropriate level
- Accommodate variation in treatment effects
- More efficient inference of regression parameters
- Using all the data to perform inferences for groups with small sample size
- Predictions fully accounting for uncertainty and variability
- Prior information

Datasets are stochastic realisations of a process

```
data1=rnorm(100, 2 + 1.6*x, 5)
```


```
data2=rnorm(100, 2 + 1.6*x, 5)
```

Histogram of data1



data1

Histogram of data2




data2

These two datasets are different, even though they arise from same process.

Data inform about process, but have 'noise' too.

Hierarchical Bayes: data, process, parameters



Clark et al. 2006, Clark 2007


$$\begin{aligned} &f(\text{data}, \text{process}, \text{parameters}) \\ &\propto f(\text{data} | \text{process}, \text{parameters}) \\ &\quad \times f(\text{process} | \text{parameters}) \\ &\quad \times f(\text{parameters}). \end{aligned}$$

Exercise

Does sex influence height?

Bayesian logistic regression

Relationship between tree size and mortality



```
logreg <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(dead)){  
    dead[i] ~ dbern(pdeath[i])  
    logit(pdeath[i]) <- mu + beta*dbhc[i]  
  }  
  
  # PRIORS  
  mu ~ dnorm(0, .001)  
  beta ~ dnorm(0, .001)  
}
```

```
data <- list(dead = trees$dead,  
             dbhc = trees$dbh.c)  
m5 <- jags(data,  
           model.file = logreg,  
           parameters.to.save = c("mu", "beta"),  
           n.chains = 3,  
           inits = NULL,  
           n.iter = 4000,  
           n.burnin = 2000)
```


Results

Inference for Bugs model at "C:/Users/FRS/AppData/Local/Temp/RtmpisLmp5
3 chains, each with 4000 iterations (first 2000 discarded), n.thin = 2
n.sims = 3000 iterations saved

	mu.vect	sd.vect	2.5%	97.5%	Rhat	n.eff
beta	0.056	0.014	0.028	0.086	1.003	880
mu	-3.484	0.250	-3.951	-3.043	1.002	1500
deviance	346.416	12.370	343.740	352.129	1.001	3000

For each parameter, n.eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor (at convergence, Rhat=

DIC info (using the rule, $pD = \text{var}(\text{deviance})/2$)

$pD = 76.6$ and $DIC = 423.0$

DIC is an estimate of expected predictive error (lower deviance is better)

```
logreg <- glm(dead ~ dbh.c, data = trees, family = binomial)
```

```
glm(formula = dead ~ dbh.c, family = binomial, data = trees)
```

```
      coef.est coef.se
```

```
(Intercept) -3.44      0.21
```

```
dbh.c        0.05      0.01
```

```
---
```

```
  n = 1000, k = 2
```

```
residual deviance = 343.7, null deviance = 360.9 (difference = 17.2)
```

STAN

STAN models are defined similarly to JAGS
(but see STAN manual for differences)

Our Bayesian linear regression in JAGS

```
bayes.lm <- function(){  
  
  # LIKELIHOOD  
  for (i in 1:length(height)){  
    height[i] ~ dnorm(mu[i], tau)      # tau = precision (inverse of vari  
    mu[i] <- alpha + beta*dbh[i]      # centred diameter  
  }  
  
  # PRIORS (vague or weakly informative)  
  alpha ~ dunif(1, 100)      # prior for average height of a 25-cm-DBH t  
  beta ~ dunif(0, 10)       # how much do we expect height to scale wit  
  tau <- 1/(sigma*sigma)    # tau = 1/sigma2  
  sigma ~ dunif(0, 50)     # residual standard deviation  
}
```

Bayesian linear regression in STAN

```
bayes.lm.stan <- "  
  
  data {  
    int<lower=0> N;  
    vector[N] dbhc;  
    vector[N] height;  
  }  
  
  parameters {  
    real<lower=0, upper=100> alpha;  
    real<lower=0, upper=10> beta;  
    real<lower=0, upper=50> sigma;  
  }  
  
  model {  
    height ~ normal(alpha + beta * dbhc, sigma);  
  }  
"
```

```
library(rstan)
lm.stan <- stan(model_code = bayes.lm.stan,
               data = list(height = trees$height, dbhc = trees$dbh.c,
                           N = nrow(trees)),
               pars = c("alpha", "beta", "sigma"),
               chains = 3,
               iter = 2000,
               init = "random")
```

Similar results as JAGS

Inference for Stan model: 57b6d06b34f74072ecc985613f1ce936.
3 chains, each with iter=2000; warmup=1000; thin=1;
post-warmup draws per chain=1000, total post-warmup draws=3000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
alpha	34.73	0.00	0.13	34.48	34.64	34.73	34.81	34.9
beta	0.62	0.00	0.01	0.60	0.61	0.62	0.62	0.6
sigma	4.10	0.00	0.09	3.93	4.03	4.10	4.16	4.2
lp__	-1905.70	0.03	1.19	-1908.75	-1906.23	-1905.36	-1904.84	-1904.4

	n_eff	Rhat
alpha	2671	1
beta	2462	1
sigma	2334	1
lp__	1526	1

Samples were drawn using NUTS(diag_e) at Tue Nov 15 17:29:09 2016.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

<https://github.com/stan-dev/example-models/wiki>

**To fit 'standard' Bayesian GLMs and
GLMMs, use rstanarm**

RStanArm

Bayesian applied regression modeling via Stan

RStanArm allows users to specify models via the customary R commands, where

- models are specified with formula syntax,
- data is provided as a data frame, and
- additional arguments are available to specify priors.

Bayesian modelling in one line of code!

```
library(rstanarm)
easybayes <- stan_glm(height ~ dbh.c, data = trees, family = gaussian,
                      prior_intercept = normal(30, 30),
                      prior = normal(0, 10))
```

and results make sense :)

```
stan_glm(formula = height ~ dbh.c, family = gaussian, data = trees,  
prior = normal(0, 10), prior_intercept = normal(30, 30))
```

Estimates:

	Median	MAD_SD
(Intercept)	34.7	0.1
dbh.c	0.6	0.0
sigma	4.1	0.1

Sample avg. posterior predictive
distribution of y (X = xbar):

	Median	MAD_SD
mean_PPD	36.5	0.2

Observations: 1000 Number of unconstrained parameters: 3

Less-tech version

More-tech version

Everything

R Users Will Now Inevitably Become Bayesians

26

There are several reasons why everyone isn't using Bayesian methods for regression modeling. One reason is that Bayesian modeling requires more thought: you need pesky things like priors, and you can't assume that if a procedure runs without throwing an error that the answers are valid. A second reason is that MCMC sampling — the bedrock of practical Bayesian modeling — can be slow compared to closed-form or MLE procedures. A third reason is that existing Bayesian solutions have either been highly-specialized (and thus inflexible), or have required knowing how to use a generalized tool like BUGS, JAGS, or Stan. This third reason has recently been shattered in the R world by not one but two packages: `brms` and `rstanarm`. Interestingly, both of these packages are elegant front ends to Stan, via `rstan` and `shinystan`.



Slides and source code available at

<https://github.com/Pakillo/Bayes-R-JAGS-intro>